

## NAME

a.s - details of CO-IDRIS startup

doshdr.86 gives control to IDRIS at line 9 via a call to \_main. \_main records the segment registers, sets up a stack, and calls main with a pointer to the command line. On return, \_main clears the dynamic data area and the user area memory, lines 18 to 30. \_main calls init to start processes 0 and 1. init places the code between lines 108 and 136 into user memory and executes it as process 1. At some point, the exec system call (line 126) or the /odd/init program will roadblock.

At that point control will resume at line 33 with the resident running as process 0. swapp is the process 0 code; it never returns.

exit (lines 37 to 41) is provided for the C library. The quit command, via the undo code, calls exit to return to DOS.

vectab, (lines 48 to 54), is a table of int locations to be plugged at startup (see the main.c manual page in this appendix). The entries are three words in length: the first word is the interrupt number to plug (the location modified is the int number \* 4); the second word is the address of the routine to be entered (ip value); and the third word is the code section to run in at the time of the interrupt (cs value). Section -1 is taken to mean the resident data section. All interrupts run in the data section at first. The -1 value at line 54 is the end of the table.

The first entry, line 49, is the clock interrupt. It is initialized to int 0 in order to disable it. The startup code in main may patch this to another value, in order to enable the clock. For this reason, the clock entry must be first in the table. See the section below on Adding a Clock for instructions on implementing your own hardware clock when not running under DOS.

Line 50 is the DOS system call intercept, and must be the second entry. If you are making an IDRIS/86 independent of DOS you will have to change line 49 and delete line 50.

Lines 51 to 53 plug the hardware exception vectors.

User programs enter the resident at sysent (line 63). Every user program reserves the first 32 bytes of its text space for system entry code. The system entry code is lines 111 to 122. The instruction at line 12 sets up this code. The exec system call checks that the first 4 bytes of the user program matches the 4 bytes at usrsvc, lines 111 to 114. If this is the case, the next 28 bytes (lines 115 to 121) are moved to the user space.

Hardware exceptions overflow, trace trap, and zero divide are dealt with in lines 78 to 94.

The csw (get console switches) system call (lines 96 to 100) is redefined for CO-IDRIS to cause a return to DOS.

The IBM serial port interrupts enter at `aca0` and `aca1`. The `int` instructions were plugged by a call to `vector` in `acaop` (open) in the `aca.c` driver file. This interrupt code, lines 139 to 157, should be used as a model for adding other interrupting devices. The hardware interrupt passes control, via the proper `int` location, to `aca0` (line 141). The call to `trapint`, line 142, does several things: saves registers, enters system mode, checks whether the user or the resident was interrupted, sets up a stack, sets `ax` to the return location, sets `dx` to the return segment, sets `cx` to 0 if the system was running, or 0x100 if the user was running, and sets up for a proper return from the interrupt. This interrupt code passes an argument to the C routine in the driver. `aca0` pushes a 0 onto the stack (lines 143 to 145); `aca1` pushes a 1 onto the stack (lines 156, 157, and 145). `ax` is the return location in the resident text space. `ax` is decremented once for each argument pushed; this causes a pop for each argument on return. Up to 3 arguments (and 3 dec `ax`'s) may be performed. The C routine is entered via a return intersegment (lines 148 to 151). Remember: interrupts are handled in the data section, but the C code is handled in the text section.

#### Adding a Clock

CO-IDRIS comes with special code to deal with clock interrupts. This code is entered at `tick`; the hook is made in line 49. To make a clock for an IDRIS running independently of DOS, change line 49 to:

```
.word 0xVV, mytick, -1

and add this code to a.s:
```

```
mytick:
    call trapint
    push ax
    push dx
    lea ax, _clock
    push ax
    reti
```

0xVV is the clock interrupt number. The above code is the model for a device interrupt with no arguments passed. Note: consult `main.c` for setting the proper value of `tinc`, the clock frequency indicator.

If you need to adjust the clock hardware, do so just before the `reti` instruction. Registers `ax`, `dx` and `cx` are available. Preserve all other registers.

## Appendix

a.s code

## NAME

a.s code - starting up CO-IDRIS

```

1: / CO-IDRIS STARTUP
2: /  copyright (c) 1984 by Whitesmiths, Ltd.
3:
4: //////////////////////////////////
5: /  BEGINNING OF PROGRAM  /
6: //////////////////////////////////
7: .text
8:     .public __main      / doshdr gives control here
9: __main:
10:     mov     __cseg,cs    / write down resident seg register
11:     mov     __dseg,ds
12:     mov     syseg,ds     / system call entry seg
13:     pop     cx
14:     pop     cx           / command line ptr
15:     lea     sp,_u0       / init stack
16:     push    cx           / pass command line ptr
17:     call    __main       / do command line
18:     .cseg
19:     mov     ax,#2        / DOS top of memory
20:     mov     __topseg,ax
21:     lea     di,__edata   / clear memory
22:     mov     ax,ds
23:     mov     es,ax
24:     mov     cx,_systop
25:     sub     cx,di
26:     shr     cx
27:     sub     ax,ax
28:     cld
29:     rep
30:     stos.w [di]
31:     cli
32:     call    __init       / start procs 0 and 1
33:     call    __swapp
34:     jmp.s   .            / should never return
35:
36:     .public __exit
37: __exit:
38:     mov     ax,0x4c      / dos 2.2 exit
39:     push    ax
40:     call    __dosy
41:     jmp.s   .
42:
43: //////////////////////////////////
44: /  INTERFACE CODE  /
45: //////////////////////////////////
46: .data
47:     .public __vectab
48: __vectab:
49:     .word   0x00,tick,-1  / MUST BE FIRST, shdb 0x08 if used
50:     .word   0x21,DOSc,-1 / MUST BE SECOND, DOS call intercept

```

```

51:      .word    0x00,zdiv,-1
52:      .word    0x01,trace,-1
53:      .word    0x04,oflo,-1
54:      .word    -1
55:
56:      SIGDOM = 7
57:      SIGHNG = 6
58:      SIGSVC = 16
59:      SIGTRC = 5
60:
61:      /   supervisor call
62:      .public sysent
63:      sysent:
64:      cli
65:      svc:
66:      call     trapint
67:      add      cx,SIGSVC
68:      trap1:
69:      sti
70:      push     cx
71:      dec      ax
72:      push     ax
73:      push     dx
74:      lea      ax,_trap
75:      push     ax
76:      reti
77:
78:      /   overflow trap
79:      oflo:
80:      call     trapint
81:      add      cx,SIGHNG
82:      jmp.s    trap1
83:
84:      /   trace trap
85:      trace:
86:      call     trapint
87:      add      cx,SIGTRC
88:      jmp.s    trap1
89:
90:      /   zero divide trap
91:      zdiv:
92:      call     trapint
93:      add      cx,SIGDOM
94:      jmp.s    trap1
95:
96:      /   undo interrupts and shutdown
97:      .public _csw
98:      _csw:
99:      call     _undo
100:      ret
101:
102:      //////////////////////////////////
103:      /   PROCESS ONE   /
104:      //////////////////////////////////

```

## Appendix

a.s code

```

105: .data
106: .public _p1boot, _p1end
107: .even
108: _p1boot:
109:
110: .public __usrsvc, syseg / USER MODE SYSTEM CALL
111: __usrsvc:
112: jmp.s 2f / this area re-written by resident
113: .word 0x6969
114: 1:
115: pop cx / IDRIS system call (prototype)
116: pushf
117: push cs
118: push cx
119: jmp.i sysent,*0
120: syseg = .-2
121: . = __usrsvc+32 / first 32 bytes controlled by
/ resident
122: 2:
123: push p1vec+2
124: push p1vec
125: push cx / dummy return link
126: mov ax,11 / exec
127: call 1b
128: mov ax,38 / undo()
129: call 1b
130: jmp.s .
131: p1vec = . - _p1boot
132: .word p1path, p1vec+4, p1path+5, 0
133: p1path = . - _p1boot
134: "/odd/init\0"
135: .even
136: _p1end:
137:
138: / IBM Asynchronous Communications Adaptor (com1, com2)
139: .data
140: .public _aca0
141: _aca0:
142: call trapint
143: mov cx,0 / acaint(0)
144: 1:
145: push cx / push line number
146: dec ax / do extra pop on return for arg passed
147: push ax / return loc for acaint()
148: push dx / cseg for reti
149: lea ax,_acaint
150: push ax / offset for reti
151: reti
152:
153: .public _aca1
154: _aca1:
155: call trapint
156: mov cx,1 / acaint(1)
157: jmp.s 1b

```

## NAME

aca.c - IBM PC-XT serial line driver

ACAVEC (lines 11 to 13) is the structure where the adaptor I/O registers are located. These registers are, in order of successive port addresses: the character buffer (cbuf), the interrupt enable register (ier), the interrupt request register (iir), the line control register (lcr), the modem control register (mcr), the line status register (lsr), and the modem status register (msr).

ACASPEED (lines 15 to 17), is the structure containing the 2-byte speed code register.

The manifest constants, lines 27 to 89, define the I/O and interrupt controller register contents.

The number of adaptors and their addresses is defined in lines 95 to 96. nmaca (line 111) defines the names of the adaptor /dev entries. Note: this driver can support 2 physical lines, with a primary and alternate minor number for each. The primary number, com1 or com2, is used by and the alternate number, com1.lnk or com2.lnk, is used by cu to take control of the line away from paca (line 107) is a pointer to the line controller currently in effect. INTMASK (line 100) determines the physical interrupt given the minor number; and ADDRMASK (line 101) determines the physical line number given the minor number. The driver is coded to handle the case where there may be multiple adaptors sharing the same interrupt.

The resident knows the entry points to the driver via acaudev (line 105) which is referenced in main.c.

acastab (lines 117 to 118) is the table of divisor values for each baud rate.

acacl (lines 120 to 147) is called by the resident for every close of an adaptor line. Nothing happens unless this is the last close (lines 135 to 136). The user program roadblocks until output is flushed (line 137). The line controller is reset (lines 138 to 140). If the line controller being closed is currently being used, its alternate is put in use. If the alternate is not open either, the interrupts for the line are turned off; otherwise the baudrate of the alternate line is used. Note: in this context, com1 and com1.lnk are alternates of one another, as are com2 and com2.lnk.

acaop (lines 155 to 228) is called by the resident for each open system call (/dev/com1, et. al.). devm, line 166, is the minor number (0-3); whereas mdevm, line 169, is the physical line number (0-1).

acaop ensures that only defined minor numbers can be opened (lines 171 to 175). pv, line 176, gets the I/O address of the physical line. pt (line 177) goes to the address of the line controller to get the minor number.

On the first open to a line (lines 178 to 206), the interrupt vector is plugged and the default line speed and stty-settable parameters are initialized. dseg, (lines 160, 183, and 187) is the resident data section

(ds register). `aca0` and `aca1` are in module `a.s.` `aca0` and `aca1` call `acaint` at line 290.

If the alternate line controller is open, the speed and parameters are copied (lines 197 to 198). If the line is not opened to either line controller (line 209) an attempt is made to set up the hardware for operation (lines 211 to 213). If the hardware is installed (lines 214 to 217) the pointer to the line controller is established, otherwise a failure is reported. If the line is already opened, an `open` to `comm1.lnk` takes precedence over `com1`, and `com2.lnk` takes precedence over `com2` (lines 219 to 220).

The line is marked "open" if the hardware appears to be installed (lines 221 to 226).

`acard`, lines 232 to 238, is called by the resident for every read system call. `ttread` does the work, given a pointer to the line controller for the minor number. Note: the program will roadblock in `ttread` until there is input. If the line controller is not current, no input will appear.

`acaset`, lines 242 to 262, sets baud rate and configuration. Note: speed code 0 means that modem signals are turned off (lines 251 to 252).

`acasz`, lines 266 to 276, is called by the resident for each `stty` and `gtty` system call. `ttset` does all the work unless the call is an `stty`. For an `stty`, `getfl` is `NO`, and the driver must set the baudrate.

`acawr` (lines 280 to 286) is called by the resident for each write system call. `ttwrite` queues the characters and calls `acago`. Note: the program will roadblock here when the output queue backs up (reaches `out_hi`).

`acaint` (lines 290 to 317) is entered from the interrupt code in `a.s.` The interrupt routines pass the adaptor number as an argument (lines 290 to 291). The `for` loop, line 300, really only executes once. The code was constructed to allow for multiple physical lines sharing the same interrupt vector. Nothing is done unless there really is an interrupt request (line 301). The status registers are read at lines 303 to 304. If there is input (line 305) the character is read at line 307. If an overrun is indicated (line 308) a `DEL` code is inserted to indicate that data has been lost. This is necessary for `cu` to function properly. If a framing error or break is indicated (lines 310 to 311) the character is treated as a `DEL` code. The transmit-go routine, `acago`, is called if there is an output interrupt. The interrupt controller chip is cleared on the way out (line 317). Note: `ttin` will result in a call to `acago` if echoing is being done.

`acago` (lines 322 to 360) is called by `ttwrite` and `ttin` to start the next character transmission. A pointer to `acago` is established by `acaop` (line 196). `acago` first checks whether the transmitter is busy (line 332). `T_STOP` is set if an `X-OFF` (`ctl-s`) was received. `T_TIMER` is set if a delay is in progress (lines 354 to 355). `deqc` returns the next character (or -1 if none is present) at line 333. In raw mode, 8 bits are transmitted as is (lines 335 to 336). In normal mode, 7 bits and parity is sent (lines 337 to 351). `M_EVEN` and `M_ODD` determine 1 parity (line 341 to 343), even

parity (lines 344 to 346), odd parity (lines 347 to 348), or 0 parity (default). If the queue is empty or the low-water mark is reached, processes blocked on output are awakened (lines 358 to 359).



## Appendix

aca.c code

## NAME

```

aca.c code - asynchronous communication adapter
1:  /* IBM PC ASYNCHRONOUS COMMUNICATIONS ADAPTER
2:  *   copyright (c) 1983 by Whitesmiths, Ltd.
3:  */
4:  #include <std.h>
5:  #include <res.h>
6:  #include <bio.h>
7:  #include <cio.h>
8:
9:  /* aca registers
10:  */
11:  typedef struct {
12:      UTINY obuf, ier, iir, lcr, mcr, lsr, msr;
13:  } ACAVEC;
14:
15:  typedef struct {
16:      UTINY lsbspeed, msbspeed;
17:  } ACASPEED;
18:
19:  /* onunit control struct
20:  */
21:  typedef struct {
22:      TEXT *next;
23:      VOID (*on_fn)();
24:      BYTES on_a1, on_a2;
25:  } ONUNIT;
26:
27:  /* 8259 interrupt controller
28:  */
29:  #define INTA00 0x20 /* control port */
30:  #define INTA01 0x21 /* int mask port */
31:  #define EOI 0x20 /* end-of-interrupt */
32:  #define IRQ0 1 /* int mask bits */
33:  #define IRQ1 2
34:  #define IRQ2 4
35:  #define IRQ3 8
36:  #define IRQ4 16
37:  #define IRQ5 32
38:  #define IRQ6 64
39:  #define IRQ7 128
40:
41:  /* Line-Control Register
42:  */
43:  #define WLS0 1 /* word length select bit 0 */
44:  #define WLS1 2 /* word length select bit 1 */
45:  #define STB 4 /* number of stop bits */
46:  #define PEN 8 /* parity enable */
47:  #define EPS 16 /* even parity select */
48:  #define STICK 32 /* stick parity */
49:  #define SBREAK 64 /* set break */
50:  #define DLAB 128 /* divisor latch access bit */
51:
52:  #define CONFIG (WLS1|WLS0) /* 8-bits, no parity */

```

```

53:
54:  /* Modem Control Register
55:  */
56:  #define DTR      1      /* data terminal ready */
57:  #define RTS      2      /* request to send */
58:  #define OUT1     4      /* output 1 */
59:  #define OUT2     8      /* output 2 */
60:  #define LOOP     16     /* loopback mode */
61:
62:  #define LINEON (OUT2|DTR|RTS) /* enable ints and signals */
63:  #define LINEOFF OUT2         /* enable ints, no signals */
64:
65:  /* Modem Status Register
66:  */
67:  #define DCTS      1      /* delta clear to send */
68:  #define DDSR      2      /* delta data set ready */
69:  #define TERI      4      /* trailing edge ring indicator */
70:  #define DRLSD     8      /* delta receive line signal detect */
71:  #define CTS       16     /* clear to send */
72:  #define DSR       32     /* data set ready */
73:  #define RI        64     /* ring indicator */
74:  #define RLSD     128     /* receive line signal detect */
75:
76:  /* Interrupt Enable Register
77:  */
78:  #define ENABLE 0x7      /* modem status|recv status|xmit|recv */
79:  #define DISABLE 0
80:
81:  /* Line Status Register
82:  */
83:  #define DR        1      /* data ready */
84:  #define OR        2      /* overrun error */
85:  #define PE        4      /* parity error */
86:  #define FE        8      /* framing error */
87:  #define BI        16     /* break interrupt */
88:  #define THRE      32     /* transmitter holding register empty */
89:  #define TSRE      64     /* transmitter shift register empty */
90:
91:  /* define aca addresses and number
92:  * NOTE: You make entries here, in acaop(), and in \
93:  * a.s in order
94:  * to add more aca lines.
95:  */
95:  LOCAL ACAVEC *acaaddr[] {0x3f8, 0x2f8};
96:  #define NACA (sizeof (acaaddr) / sizeof (acaaddr[0]))
97:
98:  /* minor number encoding
99:  */
100:  #define INTMASK 1
101:  #define ADDRMASK 1
102:
103:  /* I/O control structure
104:  */
105:  CDEVSW acadev {&acaop, &acacl, &acard, &acawr, \

```

## Appendix

aca.c code

```

        &acasg, nmaca);
106:  LOCAL TTY aca[NACA*2] {0};
107:  LOCAL TTY *paca[NACA] {0};
108:
109:  /* device name displayed by /dev/cnames
110:  */
111:  LOCAL TEXT nmaca[] {"com1 com2 com1.lnk com2.lnk"};
112:
113:  /* Speed codes (divisor latch settings)
114:  * 0, 50, 75, 110, 134.5, 150, 200, 300,
115:  * 600, 1200, 1800, 2400, 4800, 9600, exta, extb
116:  */
117:  LOCAL UCOUNT acastab[16] {0, 2304, 1536, 1047, 857, \
        768, 576, 384,
118:        192, 96, 64, 48, 24, 12, 0, 0};
119:
120:  /* close comm line
121:  */
122:  LOCAL VOID acacl(dev)
123:  DEV dev;
124:  {
125:  IMPORT ACAVEC #acaaddr[];
126:  IMPORT BITS ttypes;
127:  IMPORT TTY aca[];
128:  IMPORT TTY *paca[];
129:  FAST COUNT devm = dminor(dev);
130:  FAST ACAVEC #pv = acaaddr[devm & ADDRMASK];
131:  FAST TTY #pt = &aca[devm];
132:  BITS ps;
133:  TTY #pa = &aca[devm < NACA ? devm + NACA : devm - NACA];
134:
135:  if (--pt->t_open)
136:  return;
137:  wflush(pt);
138:  ps = spl(ttypes);
139:  pt->t_stat = 0;
140:  settyp(dev, NO);
141:  if (pt == paca[devm & ADDRMASK])
142:  if (!(paca[devm & ADDRMASK] = pa->t_stat ? pa : NULL))
143:  out(&pv->ier, DISABLE);
144:  else
145:  acaset(pa);
146:  spl(ps);
147:  }
148:
149:  /* open comm line
150:  * NOTE: You have to make entries in the switch below to add
151:  * more aca lines. The first argument to vector() is \
152:  * the interrupt
153:  * vector number; the second arg is the interrupt \
154:  * routine address
155:  * in a.s; the third arg is the resident DS register.
156:  */
157:  LOCAL VOID acaop(dev)

```

```

156:     DEV dev;
157:     {
158:     IMPORT ACAVEC #acaaddr[];
159:     IMPORT BITS ttypts;
160:     IMPORT BYTES dseg;
161:     IMPORT TTY aca[];
162:     IMPORT TTY #paca[];
163:     IMPORT VOID acago();
164:     IMPORT VOID aca0(), aca1();
165:     FAST ACAVEC #pv;
166:     FAST COUNT devm = dminor(dev);
167:     FAST TTY #pt;
168:     BITS ps;
169:     COUNT mdevm = devm & ADDRMASK;
170:
171:     if (NACA*2 <= devm)
172:     {
173:         uerror(ENXIO);
174:         return;
175:     }
176:     pv = acaaddr[mdevm];
177:     pt = &aca[devm];
178:     if (!pt->t_go)
179:     {
180:         switch (devm & INTMASK)
181:         {
182:         case 0:
183:             vector(12, &aca0, dseg); /* /dev/com1,com1.lnk */
184:             out(INTA01, in(INTA01) & ~IRQ4);
185:             break;
186:         case 1:
187:             vector(11, &aca1, dseg); /* /dev/com2,com2.lnk */
188:             out(INTA01, in(INTA01) & ~IRQ3);
189:             break;
190:         /* add entries here */
191:         default:
192:             uerror(ENXIO);
193:             return;
194:         }
195:         pt->t_dev = dev;
196:         pt->t_go = &acago;
197:         if (paca[mdevm])
198:             movbuf(&paca[mdevm]->t_speeds, &pt->t_speeds, 6);
199:         else
200:         {
201:             pt->t_speeds = 0x0909; /* 1200 baud */
202:             pt->t_erase = CERASE;
203:             pt->t_kill = CKILL;
204:             pt->t_flag = M_XTABS|M_ECHO|M_CRTLF;
205:         }
206:     }
207:     settyp(dev, YES);
208:     ps = spl(ttypts);
209:     if (!paca[mdevm])

```

## Appendix

aca.c code

```

210:         {
211:             out(&pv->lcr, CONFIG);
212:             acaset(pt);
213:             out(&pv->ier, ENABLE);
214:             if (in(&pv->lcr) != CONFIG || in(&pv->ier) != ENABLE)
215:                 uerror(ENODEV);
216:             else
217:                 paca[mdevm] = pt;
218:         }
219:     else if (NACA <= devm)
220:         paca[mdevm] = pt;
221:     if (!uerror(0))
222:     {
223:         if (!(pt->t_stat & T_OPEN))
224:             pt->t_stat = T_OPEN | T_CARR;
225:         ++pt->t_open;
226:     }
227:     spl(ps);
228: }
229:
230: /* read comm line
231: */
232: LOCAL VOID acard(dev)
233:     DEV dev;
234: {
235:     IMPORT TTY aca[];
236:
237:     ttread(&aca[dminor(dev)]);
238: }
239:
240: /* set aca speed
241: */
242: LOCAL acaset(pt)
243:     TTY *pt;
244: {
245:     IMPORT ACAVEC *acaaddr[];
246:     IMPORT UCOUNT acastab[];
247:     FAST ACAVEC *pv = acaaddr[dminor(pt->t_dev) & ADDRMASK];
248:     FAST UCOUNT speed;
249:     FAST UTINY lcrreg;
250:
251:     if (!(speed = acastab[pt->t_speeds & 0xf]))
252:         out(&pv->mcr, LINEOFF);
253:     else
254:     {
255:         lcrreg = in(&pv->lcr);
256:         out(&pv->lcr, lcrreg | DLAB);
257:         out(&pv->lsbspeed, speed);
258:         out(&pv->msbspeed, speed >> 8);
259:         out(&pv->lcr, lcrreg);
260:         out(&pv->mcr, LINEON);
261:     }
262: }
263:

```

```

264:  /* set or get tty mode
265:  */
266:  LOCAL VOID acasg(dev, getfl)
267:      DEV dev;
268:      COUNT getfl;
269:      {
270:      IMPORT TTY aca[];
271:      IMPORT TTY *paca[];
272:      FAST TTY *pt = &aca[dminor(dev)];
273:
274:      if (!ttset(pt, getfl) && pt == paca[dminor(dev) & ADDRMASK])
275:          acaset(pt);
276:      }
277:
278:  /* write comm line
279:  */
280:  LOCAL VOID acawr(dev)
281:      DEV dev;
282:      {
283:      IMPORT TTY aca[];
284:
285:      ttwrite(&aca[dminor(dev)]);
286:      }
287:
288:  /* adapter interrupt
289:  */
290:  VOID acaint(devm)
291:      UTINY devm;
292:      {
293:      IMPORT ACAVEC *acaaddr[];
294:      IMPORT TTY aca[];
295:      IMPORT TTY *paca[];
296:      FAST ACAVEC *pv;
297:      FAST TTY *pt;
298:      UTINY iir, lsr, msr, c;
299:
300:      for (; devm < NACA; devm += 2)
301:          if ((pt = paca[devm]) && !((iir = in(&(pv = \
302:              acaaddr[devm])->iir)) & 1))
303:              {
304:              lsr = in(&pv->lsr);
305:              msr = in(&pv->msr);
306:              if (lsr & DR)
307:                  {
308:                  c = in(&pv->cbuf);
309:                  if (lsr & OR)
310:                      ttin(CINTR, pt);
311:                  if (lsr & (FE|BI))
312:                      c = CINTR;
313:                  ttin(c, pt);
314:                  }
315:              if (((iir & 7) == 2) && (lsr & THRE))
316:                  acago(pt);
317:              }

```

```

317:         out(INTA00, EOI);
318:     }
319:
320:     /* start transmitter
321:     */
322:     LOCAL VOID acago(pt)
323:         FAST TTY *pt;
324:     {
325:         IMPORT ACAVEC *acaaddr[];
326:         IMPORT COUNT out_lo;
327:         IMPORT UTINY cmaptab[];
328:         IMPORT VOID ttrstart();
329:         FAST ACAVEC *pv = acaaddr[dminor(pt->t_dev) & ADDRMASK];
330:         FAST COUNT c;
331:
332:         if (!(in(&pv->lsrc) & THRE) || pt->t_stat & \
333:             (T_STOP|T_TIMER) ||
334:             (c = deqc(&pt->t_outq)) < 0)
335:             return;
336:         else if (pt->t_flag & M_RAW)
337:             out(&pv->cbuf, c);
338:         else if (c < 0200)
339:         {
340:             switch (pt->t_flag & (M_EVEN|M_ODD))
341:             {
342:             case M_EVEN|M_ODD:
343:                 c = 0200;
344:                 break;
345:             case M_EVEN:
346:                 c = cmaptab[c] & 0200;
347:                 break;
348:             case M_ODD:
349:                 c = cmaptab[c] & 0200 ^ 0200;
350:             }
351:             out(&pv->cbuf, c);
352:         }
353:         else
354:         {
355:             timeout(&ttrstart, pt, c & 0177);
356:             pt->t_stat |= T_TIMER;
357:             return;
358:         }
359:         if (pt->t_outq.c_num == 0 || pt->t_outq.c_num == out_lo)
360:             wakeup(&pt->t_outq);

```

## NAME

bio.h - driver header file used by all drivers

Each 512-byte buffer is controlled by a buffer controller, which is of type BUF (lines 20 to 32). Buffer controllers are maintained as a doubly-linked list when on the free list (f\_next and f\_prev) or in the cache (b\_next and b\_prev), and as a null-terminated list when enqueued for a device (d\_next and d\_prev). Each device driver has a DEVTAB, lines 36 to 44, which heads up its chain. The codes for b\_flag, lines 5 to 18, indicate the state of a buffer, and communicate between the driver and the resident.

B\_ASYNC is set for read-ahead and write-behind operations.

B\_BUSY is set when I/O is in progress.

B\_DIRTY is set when a buffer contents have changed, but have not been written to their block device.

B\_ERROR is set by the driver when the I/O has failed.

B\_READ is set before calling the driver to indicate a read operation, as opposed to a write.

B\_TAPE is set by the device driver for blocked tape I/O. Its purpose is to defeat read-ahead and write-behind to keep tape blocks in order.

B\_VALID is set by a call to iodone; it indicates that the contents of a particular buffer are valid.

B\_WANT is set by processes waiting on a busy buffer; it causes a wakeup to be done when the I/O completes.

B\_CHR is set by physio to mark a character special (raw mode) transfer, as opposed to a block special I/O.

B\_CTRL is not really a buffer flag, but a parameter sent to physio to indicate a control operation (as opposed to a data transfer). It is used to defeat buffer count and address checks, and to keep the requesting process from being locked in memory.

B\_MAP and B\_RESRC are used in the PDP-11 and the VAX to control mapping hardware. B\_MAP is set for all time for buffers in the resident address space. B\_RESRC is set when some mapping registers are allocated.

BDEVSW (lines 46 to 54) defines the table of entry points published by each block device driver.

CDEVSW (lines 58 to 65) defines the table of entry points published by each character device driver.



## NAME

bio.h code - blocked I/O header file

```

1:  /* HEADER FOR BLOCKED I/O OPERATIONS
2:      *  copyright (c) 1979 by Whitesmiths, Ltd.
3:      */
4:
5:      /* codes for b_flag
6:      */
7:      #define B_ASYNC 00001
8:      #define B_BUSY 00002
9:      #define B_DIRTY 00004
10:     #define B_ERROR 00010
11:     #define B_READ 00020
12:     #define B_TAPE 00040
13:     #define B_VALID 00100
14:     #define B_WANT 00200
15:     #define B_CHR 00400
16:     #define B_CTRL 01000
17:     #define B_MAP 02000
18:     #define B_RESRC 04000
19:
20:     struct buf {
21:         BUF *f_next;
22:         BUF *f_prev;
23:         BUF *b_next;
24:         BUF *b_prev;
25:         BITS b_flag;
26:         DEV b_dev;
27:         BYTES b_count;
28:         ULONG b_phys;
29:         TEXT *b_addr;
30:         BYTES b_resid;
31:         BLOCK b_blkno; /* ULONG b_off[2]; */
32:     };
33:
34:     /* the device control table
35:     */
36:     #define DEVTAB struct devtab
37:     struct devtab {
38:         BUF *d_next;
39:         BUF *d_prev;
40:         BUF *b_next;
41:         BUF *b_prev;
42:         TINY d_stat;
43:         TINY d_nerr;
44:     };
45:
46:     /* the block device switch table entry
47:     */
48:     typedef struct {
49:         BOOL (*d_open)();
50:         VOID (*d_close)();

```

```
51:     VOID (*d_strat)();
52:     DEVTAB *d_tab;
53:     TEXT *d_bname;
54:     } BDEVSW;
55:
56: /* the character device switch table entry
57: */
58: typedef struct {
59:     BOOL (*d_open)();
60:     VOID (*d_close)();
61:     COUNT (*d_read)();
62:     COUNT (*d_write)();
63:     VOID (*d_sgatty)();
64:     TEXT *d_cname;
65:     } CDEVSW;
```

## NAME

cio.h - driver header file used only by terminal drivers

TTY (lines 69 to 85) controls character I/O to a terminal. The information in the structure is used by both terminal device drivers and the resident. All the information for a particular terminal is in one of these structures.

Certain character codes, lines 7 to 15, are compiled-in. CEOT is the terminal end-of-file character. CDELIM is an internal code. CERASE is the default character-delete character. CGO is the transmit-go character X-ON (ctl-q). CINTR is the keyboard interrupt character (DEL key). CKILL is the default line-delete character. CQUIT is the keyboard abort character (ctl-^). CSTOP is the transmit-stop character X-OFF (ctl-s). CALERT is the audible tone (bell).

t\_flag bits, lines 19 to 26, select the mode of the terminal. M\_RARE is set for rare mode. M\_XTABS controls tab expansion. M\_CASE controls upper-case to lower-case mapping. M\_ECHO controls full duplex operation (i.e. echoing of input). M\_CRTLF controls carriage-return to line-feed mapping. M\_RAW is set for raw mode; it supercedes M\_RARE. M\_ODD and M\_EVEN control parity bit generation.

t\_speeds fields, lines 30 to 37, hold the terminal's speed codes and some flags. t\_speeds is accessed via stty and gtty system calls, the resident, and the terminal driver. S\_ISPEED is the mask for the input baud-rate; or the baud-rate for the line if split speeds are not possible or desirable. S\_IBREAK is set by the driver when a break key is detected. S\_ILOST is set by the driver when input has been lost (data overrun). S\_IREADY is set by ttset (invoked by the resident) for each stty if input is ready. S\_OSPEED is the output baud-rate code. S\_OBREAK is a command to the transmit routine to send a break. The transmit routine resets this flag when the break is sent. S\_ONXON is a flag to the resident to inhibit the transmission of X-ON/X-OFF (flow control) when set. S\_OREADY is set by ttset (invoked by the resident) for each stty system call when the output queue is empty.

t\_stat flags, lines 41 to 48, indicate the state of the terminal. T\_BUSY is set by ttin when an X-OFF has been sent. T\_CARR is set by the driver when the line is considered connected (data set ready = true). ttin will not operate unless this flag is true. T\_ESC is set for escape sequences (\ received). T\_OPEN is set by the driver when the line is opened. ttin will not operate unless this flag is true. T\_STOP is set when an X-OFF has been received. This flag is honored by the transmit-go routine in the driver. T\_TIMER is set by the transmit-go routine in the driver when a delay is in progress. T\_CLOS is set by some drivers to indicate that the line is closing. T\_WOPEN is set by some drivers when waiting for the line to connect (data set ready = false). When the line is connected (i.e. the terminal in question is powered-up, and/or a modem has answered a call), T\_WOPEN is reset and T\_OPEN and T\_CARR are set.

CLIST is the structure of each 16-byte character buffer (lines 52 to 56).  
CHQ is the list-head for each null-terminated list of character buffers,  
(lines 60 to 65). The character queues are maintained by deqc, enqc, and  
popc.

## Appendix

cio.h code

## NAME

cio.h code - character I/O

```
1: /* HEADER FOR CHARACTER I/O OPERATIONS
2:  * copyright (c) 1979 by Whitesmiths, Ltd.
3:  */
4:
5: /* various paramters
6:  */
7: #define CEOT      0004
8: #define CDELIM    0377
9: #define CERASE    '\b'
10: #define CGO       0021
11: #define CINTR     0177
12: #define CKILL     '\25'
13: #define CQUIT     0034
14: #define CSTOP     0023
15: #define CALERT    007
16:
17: /* codes for t_flag
18:  */
19: #define M_RARE    000001
20: #define M_XTABS   000002
21: #define M_LCASE   000004
22: #define M_ECHO    000010
23: #define M_CRTLF   000020
24: #define M_RAW     000040
25: #define M_ODD     000100
26: #define M_EVEN    000200
27:
28: /* code for t_speeds
29:  */
30: #define S_ISPEED   0x000f
31: #define S_IBREAK   0x0010
32: #define S_ILOST    0x0020
33: #define S_IREADY   0x0080
34: #define S_OSPEED   0x0f00
35: #define S_OBREAK   0x1000
36: #define S_ONXON    0x2000
37: #define S_OREADY   0x8000
38:
39: /* codes for t_stat
40:  */
41: #define T_BUSY     0001
42: #define T_CARR     0002
43: #define T_ESC      0004
44: #define T_OPEN     0010
45: #define T_STOP     0020
46: #define T_TIMER     0040
47: #define T_WCLOS    0100
48: #define T_WOPEN    0200
49:
50: /* the character buffer structure
```

```
51:  /*
52:  #define CLIST    struct clist
53:  struct clist {
54:      CLIST *c_next;
55:      TEXT c_info[16 - sizeof (BYTES)];
56:  };
57:
58:  /* the character queue structure
59:  */
60:  #define CHQ struct chq
61:  struct chq {
62:      COUNT c_num;
63:      TEXT *c_first;
64:      TEXT *c_last;
65:  };
66:
67:  /* the tty control structure
68:  */
69:  #define TTY    struct tty
70:  struct tty {
71:      CHQ t_rawq;
72:      CHQ t_outq;
73:      VOID (*t_go)();
74:      DEV t_dev;
75:      BITS t_stat;
76:      COUNT t_ndel;
77:      COUNT t_col;
78:      COUNT t_prevcol;
79:      COUNT t_open;
80:      COUNT t_nin;
81:      BITS t_speeds; /* returned by stty */
82:      TEXT t_erase; /* returned by stty */
83:      TEXT t_kill; /* returned by stty */
84:      BITS t_flag; /* returned by stty */
85:  };
```

## NAME

dos13.s - description of IBM PC hard disk I/O

dos13 (lines 6 to 26) does an int 13 call to the IBM BIOS. The register setup is documented in the IBM PC-XT Technical Reference Manual.

The arguments (line 3) are: **nblk**, the number of 512-byte sectors to transfer; **op**, the operation; **offset**, the 16-bit buffer offset from base; **sec**, the sector with 2 bits of cylinder number; **cyl**, the low 8 bits of cylinder number; **unit**, the disk unit with high-bit set; **head**, the disk head (surface) number; and **base**, the 16-bit paragraph. dos13 returns 0 or an error code (lines 20 to 22).

dos13st (lines 29 to 43) is a special form of the int 13 IBM BIOS call which returns the status in the registers to a C structure. The arguments (line 28) are: **unit**, the disk unit with high-bit set; and **stat**, the address of the structure to hold status. dos13st returns 0 or an error code, as do the **cx** and **dx** registers in the **stat** buffer (lines 36 to 41).

## NAME

code for dos13.s - IBM PC hard disk I/O

```

1: / IBM PC BIOS hard disk I/O
2: /  copyright (c) 1983 by Whitesmiths, Ltd.
3: /  BITS dos13(nblk, op, offset, sec, cyl, unit, head, base)
4:
5:      .public _dos13
6:  _dos13:
7:      push    bp
8:      mov     bp,sp
9:      push    bx
10:     push    es
11:     mov     al,[bp][4]
12:     mov     ah,[bp][6]
13:     mov     bx,[bp][8]
14:     mov     cl,[bp][10]
15:     mov     ch,[bp][12]
16:     mov     dl,[bp][14]
17:     mov     dh,[bp][16]
18:     mov     es,[bp][18]
19:     int     0x13
20:     jc      1f
21:     sub     ax,ax
22: 1:
23:     pop     es
24:     pop     bx
25:     pop     bp
26:     ret
27:
28: /  BITS dos13st(unit, &stat)
29:      .public _dos13st
30:  _dos13st:
31:      push    bp
32:      mov     bp,sp
33:      mov     dl,[bp][4]
34:      mov     ah,8
35:      int     0x13
36:      jc      1f
37:      sub     ax,ax
38: 1:
39:      mov     bp,[bp][6]
40:      mov     [bp][0],cx
41:      mov     [bp][2],dx
42:      pop     bp
43:      ret

```



## NAME

main.c - file of system configuration parameters

The data specific to CO-IDRIS (lines 7 to 16) are used only by main.c and some of the CO-IDRIS device drivers. MAXCMD (line 9) is the length of the longest command line tolerated. pname (line 16) is printed by the usage message of getflags.

The default filesystems (lines 20 to 21) are encoded as a major number in the high 8 bits, and a minor number in the low 8 bits.

If swapdev and rootdev are the same (lines 25 to 27), swapadr is equal to the size of the root filesystem, and swapsiz + swapadr is equal to the size of the disk. If you put swapping on another block device (i.e. swapdev is not the same as rootdev), then swapadr should be set to 1 and swapsiz would be set to the size of the swap device - 1.

Note: a block device can be implemented on any random access media available. It is typically a disk, but could be RAM, ROM, Dectape II, magnetic bubbles, or whatever. You could put the root file-system in ROM, pipes in a RAM disk, and user space in magnetic bubbles. Leave swapping turned off, unless RAM memory is short. It is best to allocate swapping on some fixed media, like a winchester or RAM disk, although diskettes have been used in a pinch.

blkdevs (lines 32 to 38) is the table of linkages to the block device drivers. chrdevs (lines 44 to 52) is the table of linkages to the character device drivers. To remove a device from the system, replace the device name with &nobdev (for block special devices), or &nocdev (for character special devices). Delete the object module for the driver, and remove any interrupt routine calls from a.s. It is better to keep the table entries in order because the position of an entry in blkdevs or chrdevs determines the major number of the device, which has a direct impact on the contents of the /dev directory.

The character buffering limits (lines 58 to 63) are applied on a terminal-by-terminal basis, independent of the number of terminals in the system. nclist, line 69, is the number of character buffers available for the system (each 16-byte buffer holds 14 characters). When adjusting these parameters for your system, remember the statistical nature of character I/O. You also have to tune for your mix of line speeds. A printer will consume at most (out\_hi / 14) buffers while operating: the process outputting to the printer will roadblock when out\_hi characters are buffered and it will wakeup when out\_lo characters are left. A cu packet (using cu or up) is about 200 characters. Only serial lines use these buffers; bit-mapped displays like the DOS console never tie up more than 2 buffers.

in\_max is the number of characters the resident will hold onto for input. You may wish to increase in\_max if you have a data-gathering application with a microprocessor inputting at 120cps or more. Remember - don't count the baud rate; it's the number of characters being inputted per second versus the input consumption rate of the printer that matters.

`in_xoff` is the number of characters queued on input before an X-OFF character is sent. You can disable this feature for the system by setting `in_xoff` equal to or greater than `in_max`. You can disable this feature on a line-by-line basis with an `stty` call. `in_xon` determines the lower threshold for sending X-ON, once X-OFF has been sent. It needs to be less than `in_xoff`.

`in_lim` determines the maximum type-ahead in normal mode.

Much of the data area consumed by IDRIS is dynamic and tuneable, (lines 65 to 92).

The biggest user of space is the block device buffer cache. `nbufs`, line 67, determines the number of buffers and buffer controllers. (Figure on its size being about  $512 + (24 * nbufs)$  bytes). The minimum number of buffers allowed is 2. If you have 6 buffers or less, set `rahead` (line 104) to NO. 10 to 20 buffers is a comfortable number for one or two users on a winchester disk.

`(nolist * 16)` is the number of bytes of character queue consumed. `nmap * 6` is the length of an internal table: if you see the non-fatal message MAP!, increase this number. `nheap` is the number of bytes, in addition to the init code, used for the heap area. `nheap` can be very small unless the resident is linked for separate code and data space. Then, it must be 1024 to 4096, depending on load. Processes, in-core inodes, mount entries, and open files consume the heap.

`tinc` (line 97) encodes the clock frequency. The `tinc` number that the user types (`ntinc`) is rather highly encoded (lines 175 to 180).

`svdnib` (lines 83 to 92) is a table processed by `nibble` (line 199), which actually sets aside the dynamic data area. The `mask` field determines the addressing boundary; the `ptr` field indicates the cell, (or is set to NULL), which gets the address of the area. The `m1` field indicates the pointer to the number (or is set to NULL) of items to allocate. The `m2` field (meaningful if `m1` is not NULL) points to a word holding the size in bytes of each item.

`biops` and `ttyps` should be set to 0 to disable interrupts, or set to 0x200 to enable interrupts. Interrupts can be enabled for block I/O, if no block device interrupts are used. Interrupts can be enabled for character I/O, if no serial line interrupts are used. Block and terminal I/O are independent, and can interrupt each other. The clock can interrupt both terminal and block I/O processing, but be careful if `timeout` is used.

`panic` (lines 112 to 119) should not return. Note that for CO-IDRIS it returns to DOS (line 118). A subtlety to watch for: the `sync` (line 116) may not really complete if disk devices are interrupt-driven. There may need to be a respectful delay in `idloop` before quitting.

A command line, as emitted by DOS, is parsed into the usual argument vector and count (lines 147 to 166). This code is valid only for DOS-style command lines: argument 0 (the actual command) is missing, and the first byte contains a byte count of the command line.

## Appendix

main.c

getflags (lines 168 to 173) mulls over the parsed command line, converting any values given, and possibly giving a usage message and exiting (exit in a.s).

Read-ahead gets turned off when there are only a few buffers (line 174) because the extra I/O tends to be a waste.

The clock interrupt number is passed to the table in a.s at line 182.

Device codes are derived from the names supplied (lines 183 to 193). The names used are the names stored as text strings in the block device drivers.

Initialization code is re-used as heap area, if the resident is not linked for separate code and data space (lines 194 to 198).

Remove line 196 if getflags and gblkdev are not called (these are in main-sub.o, in the library file).

## NAME

main.c code - system configuration parameters

```

1:  /* GET OPTION FROM DOS COMMAND LINE
2:  */
3:  #include <std.h>
4:  #include <res.h>
5:  #include <bio.h>
6:
7:  /* data specific to Co-Idris on DOS
8:  */
9:  #define MAXCMD 256
10:  BOOL slowclock {NO}; /* clock rate <= 4Hz */
11:  BYTES clkint {0}; /* default clock int# */
12:  TEXT *pipenm {NULL}; /* pipe device name */
13:  TEXT *rootnm {NULL}; /* root device name */
14:  TEXT *swapnm {NULL}; /* swap device name */
15:  BLOCK _rootoff {2}; /* root offset for drivers */
16:  TEXT *_pname {"idris"};
17:
18:  /* root and pipe filesystems
19:  */
20:  DEV rootdev {1<<8 | 1}; /* default is b: */
21:  DEV pipedev {1<<8 | 1}; /* must be a mounted, \
                           writeable filesystem */
22:
23:  /* swap area
24:  */
25:  DEV swapdev {NODEV}; /* no swapping */
26:  BLOCK swapadr {9000}; /* offset to swap area; CANT \
                           BE ZERO */
27:  BLOCK swapsiz {0}; /* size of swap area in blocks */
28:
29:  /* block special devices
30:  * replace unused table entries with &nobdev
31:  */
32:  BDEVSW *blkdevs[] {
33:      &nobdev, /* 0: no device */
34:      &mxpdev, /* 1: dosmx (a b c ...) */
35:      &hdbdev, /* 2: pchd (hdc0 ... hdd0) */
36:      &iebdev, /* 3: dosie (idris.0 idris.1) */
37:      &mdbdev, /* 4: md (memory disk) */
38:      };
39:  COUNT nblkdev {sizeof(blkdevs) / sizeof(blkdevs[0])};
40:
41:  /* character special devices
42:  * replace unused table entries with &nocdev
43:  */
44:  CDEVSW *chrdevs[] {
45:      &memocdev, /* 0: (ps myps cnames bnames, et. al.) */
46:      &conocdev, /* 1: dos console */
47:      &acacdev, /* 2: IBM async adaptor (com1, com2) */
48:      &mxocdev, /* 3: raw DOS disk (ra rb rc ...) */

```

```

49:      &hdcddev,          /* 4: raw IBM pchd (rhdc0 ... rhdd0) */
50:      &mdcddev,          /* 5: raw memory disk (rmd0 rmd1) */
51:      &iecddev           /* 6: raw DOS file (ridris.0 ridris.1) */
52:      };
53:  COUNT nchrdev {sizeof(chrdevs) / sizeof(chrdevs[0])};
54:
55:  /* character buffering limits
56:   * values are numbers of characters per line
57:   */
58:  UCOUNT in_max {560};    /* max input buffered in raw mode */
59:  UCOUNT in_xoff {280};   /* hiwater mark when X-OFF sent */
60:  UCOUNT in_xon {140};    /* lowater mark when X-ON sent */
61:  UCOUNT in_lim {560};    /* max characters accepted in \
        normal mode */
62:  UCOUNT out_hi {280};    /* output block hiwater mark */
63:  UCOUNT out_lo {140};    /* output unblock, lowater mark */
64:
65:  /* data space allocated at startup time
66:   */
67:  BYTES nbufs {10};        /* number of 512-byte disk buffers */
68:  BYTES nrawbs {3};        /* number of 24-byte raw-buffer \
        controllers */
69:  UCOUNT nclist {50};     /* number of 16-byte character \
        buffers */
70:  UCOUNT nmap {20};       /* number of 6-byte map elements */
71:  BYTES nheap {200};       /* number of bytes to add to \
        system heap */
72:
73:  /* nibble table for allocating space at startup
74:   */
75:  typedef struct {
76:      BYTES mask, *ptr, *m1, *m2;
77:  } NIBTAB;
78:
79:  BYTES lit1 {1};
80:  BYTES lit16 {16};
81:  BYTES lit512 {512};
82:
83:  NIBTAB svdnib[] {
84:      {01, &maplist, &nmap, &mapsiz},
85:      {01, &pbufs, &nbufs, &bufsiz},
86:      {01, NULL, &nrawbs, &bufsiz},
87:      {01, &pheap, &nheap, &lit1},
88:      {017, &eheap, NULL, NULL},
89:      {017, &pclist, &nclist, &lit16},
90:      {017, &buffers, &nbufs, &lit512},
91:      {017, &systop, NULL, NULL},
92:      {-1}};
93:
94:  /* clock tick increment
95:   * tinc = 614400 / clock-rate-Hz
96:   */
97:  ULONG tinc {614400/(1193180/65536)}; /* IBM PC Hz (~18.2) */
98:

```

```

99:  /* system flags
100:  */
101:  BITS biops      {0};      /* processing level for block I/O */
102:  BITS ttyops     {0};      /* processing level for async char I/O */
103:  BOOL fppres     {NO};     /* look for 8087 if set to YES */
104:  BOOL rahead     {YES};    /* read next block on sequential \
    file read */
105:  BOOL rootro     {NO};     /* read-only root if set to YES */
106:  BYTES stdheap   {4096};    /* default stack+heap given by exec */
107:  TEXT where[]    /* system id read by /dev/where */
108:      {"DOS Idris 6/8/84\0 Serial ####-####"};
109:
110:  /* put panic message
111:  */
112:  VOID panic(s)
113:      TEXT *s;
114:      {
115:          printf("Idris crash: %s\n", s);
116:          sync();
117:          FOREVER
118:              undo();      /* return to DOS; instead of idloop() */
119:      }
120:
121:  /* MAGIC FUNCTION.
122:  * free area between _smain() and _emain()
123:  */
124:  LOCAL VOID _smain(){}
125:
126:  /* parse command line and call getflags
127:  */
128:  BOOL main(s)
129:      FAST TEXT *s;
130:      {
131:          IMPORT BOOL fppres, rahead, rootro, silent, slowclock;
132:          IMPORT BLOCK swapadr, swapsiz, _rootoff;
133:          IMPORT BYTES clkint, nbufs, vectab[];
134:          IMPORT DEV pipedev, swapdev, rootdev;
135:          IMPORT NIBTAB svdnib[];
136:          IMPORT TEXT *_pname, *swapnm, *pipenm, *rootnm, _memory[];
137:          IMPORT ULONG tinc;
138:          IMPORT VOID _emain(), _smain();
139:          IMPORT VOID _emain1(), _smain1();
140:          FAST COUNT i;
141:          FAST TEXT *t;
142:          BYTES ntinc, pac;
143:          COUNT ac, n;
144:          TEXT *av[64], *q, qbuf[MAXCMD+1];
145:          TEXT **pav;
146:
147:          for (q = qbuf, n = *s++; 0 < n; )
148:              {
149:                  while (0 < n && iswhite(*s))
150:                      --n, ++s;
151:                  if (n <= 0)

```

```

152:         break;
153:         for (; 0 < n && !iswhite(*s); --n, ++s)
154:             *q++ = *s;
155:             *q++ = '\0';
156:         }
157:         ac = 1;
158:         av[0] = _pname;
159:         for (t = qbuf; t != q; ++t)
160:             {
161:                 av[ac++] = t;
162:                 while (*t)
163:                     ++t;
164:             }
165:         pac = ac;
166:         pav = av;
167:         ntinc = 0;
168:         getflags(&pac, &pav,
169:             "+fpp,clkint#,tinc#,+l,pipe#,root#,swapoff#,\
170:             swapsize#,swap#,s,\
171:             +readonly,nbufs#:F",
172:             &fppres, &clkint, &ntinc, &_rootoff, &pipenm, &rootnm,
173:             &swapadr, &swapsiz, &swapnm, &silent,
174:             &rootro, &nbufs);
175:         rahead = 6 <= nbufs; /* turn off read-ahead if few bufs */
176:         if (!ntinc)
177:             ;
178:         else if (0xf000 <= ntinc)
179:             tinc = 614400L / (-ntinc);
180:         else
181:             tinc = ntinc;
182:         slowclock = 153600L <= tinc;
183:         vectab[0] = clkint;
184:         if (rootnm)
185:             {
186:                 rootdev = gblkdev(rootnm, "-root <root-name>\n");
187:                 pipedev = rootdev;
188:             }
189:         if (pipenm)
190:             pipedev = gblkdev(pipenm, "-pipe <block-name>\n");
191:         if (swapsiz)
192:             swapdev = swapnm ? gblkdev(swapnm, "-swap \
193:             <block-name>\n") : rootdev;
194:         else if (swapnm)
195:             usage("-swap <block-name> -swapsize # -swapoff #\n");
196:         if (cs(-1) == ds(-1))
197:             {
198:                 free((BYTES)&_emain1 - (BYTES)&_smain1, &_smain1, 0);
199:                 free((BYTES)&_emain - (BYTES)&_smain, &_smain, 0);
200:             }
201:         nibble(&_memory, &svdnib);
202:         /* MAGIC FUNCTION. end of init code, keep in order.
203:         */

```

## Appendix

**A - 32**



## NAME

pchd.c - description of IBM PC-XT hard disk driver

This disk driver calls the IBM PC-XT ROM BIOS code, int 13, as defined in the IBM PC-XT Technical Reference Manual.

This disk has a partition table as part of its boot block. There can be from 1 to 4 partitions on the disk. There are 4 slots in the partition table, as expected. However, the IBM fdisk utility uses the slots at the end of the table. Thus, if there are 2 partitions, fdisk partition 1 winds up as slot 3, and fdisk partition 2 winds up as slot 4 (given that slots 1-4 are available). The following is a table of boot-record slots and the partition numbers they may hold:

	fdisk partition numbers			
slot 1	1			
slot 2	2	1		
slot 3	3	2	1	
slot 4	4	3	2	1

The minor numbers are encoded as shown:

UUUUUPPP	
xxxxx000	refers to the entire disk, block 0 is the boot record
xxxxx001	partition 1, like fdisk
xxxxx010	partition 2
xxxxx011	partition 3
xxxxx100	partition 4
xxxxx101	not used
xxxxx110	not used
xxxxx111	not used

The high 5 bits are the disk number: C or D disk. The entries which use partitions 1 to 4 have an offset (`_rootoff`) added so that the first 2 blocks are skipped. This is to maintain compatibility with the `/dev/[a-h]` interface to the disk.

If the low 3 bits of the minor number are 0 (corresponding to `/dev/hdc0` and `/dev/hdd0`), the entire disk is accessed, and no offsets are applied. Block 0 of `/dev/hdc0` is the boot block.

The other limitation this driver deals with is the fact that DMA operations may not cross a physical 64K boundary.

BIOSRS (line 10) is a reset command. BIOSRD (line 11) is a read command. BIOSWR (line 12) is a write command. BIOSRC (line 13) is a recalibrate command (it rewinds the disk). CYLBITS (line 14) is a mask: the high 2 bits of the sector register is really the high 2 bits of the 10-bit cylinder number.

UNISHFT (line 18) enforces the minor number encoding above. The hard disk unit number is the high 5 bits of the minor number. NHD (line 19) controls the number of hard disks supported by this driver. It can be any number, and directly impacts table sizes.) NPART (line 20) is the number

of partitions the driver controls. Partitions 1 to 4 are `fdisk` partitions, and partition 0 is the entire disk. `NDPART` (line 21) is the number of partition table entries on the disk, referred to as "slots" above.

`HDTAB` (lines 25 to 27) is the per-disk information the driver maintains. `PARTITION` (lines 29 to 33) is the layout of a partition table entry as it appears on disk and in memory. `pt_secoff` is the sector offset of the start of the partition. `pt_nsec` is the number of sectors in the partition. The sectors are 512 bytes in length.

`BOOTREC` (lines 35 to 39) is the layout of the boot record on disk. Note that a boot-record consists of a bootstrap program, 4 partition table entries, and 2 signature bytes.

`hdbdev` (line 41) is the table of entry points for block I/O. `hdodev` (line 42) is the table of entry points for character (non-buffered) I/O.

`hdtab` (line 43) is used by the resident to manage the buffer cache for this disk. `DEVTAB` is defined in the manual page on `bio.h` in this appendix. Since this driver is not interrupt-driven, I/O is synchronous, and `enq` and `deq` are not needed. Also, the strategy routine completes the I/O; there is no interrupt or start routine as in other drivers.

`hd` (line 44) is a table of per-disk information, indexed by hard disk number.

`hdp` (line 45) is the table of disk partitions managed by the driver. `hdp` is indexed by both the disk number and the in-memory partition number (0-4). `mappart` (lines 205 to 214) is always used to calculate the address of the proper entry in this table, given the minor number.

`nmhd` and `nmrhd` (lines 46 to 48) are the names of the minor numbers that access this driver. If you change `NHD` (lines 44 and 45) you will have to create or delete disk names from these two strings. Note the extra blanks between "hdc4" and "hdd0": these are significant because they skip the unused minor numbers.

`hdclose` (lines 50 to 58) counts down the open count for each disk. This open count is used by `hdopen` (lines 80 to 81) as a first time flag. Note that if the hard disk is the root, swap, or pipe disk it will always stay open. If it is not (if the root is on diskette, for instance), the open count may go to 0 and the boot block will be re-read on the next open. This behavior differs from that of DOS, which only accesses the boot-record when the system is initialized. It allows `IDRIS` to be more responsive, for good or ill, to partition table changes (via `idisk` or `fdisk`).

`hdopen` (lines 62 to 108) counts opens and does a read of the partition table when the open count is 0.

A check is made for a valid hard disk number at lines 78 to 79. Opens are counted at lines 80 to 81. Nothing else is done unless the open count is 0.

The open count is 0 in this case, so the disk information is obtained from the BIOS (lines 83 to 83). An error is reported if the BIOS call fails, and nothing else is done.

The partition table is read (lines 84 to 91). A buffer is obtained to hold the boot-record using the `getblk/brelse` machinery. If the read fails or the signature bytes are not right, the partition table is considered invalid. In this case a message is printed (line 88), the buffer is released (line 89), and an error is reported (line 90).

If all is well so far, the partition table actually used by the driver is set up (lines 92 to 108). The per-disk information needed to calculate seek addresses from block numbers is computed (lines 94 to 96). `hd_ncyl` holds the number of cylinders. `hd_nspt` holds the number of sectors per track (per head). `hd_nspc` holds the number of sectors per cylinder. "partition 0" is set up to be the size of the entire disk (line 97). The in-memory partition table is copied from the boot-record and made to line up with the fdisk numbering scheme explained above (lines 98 to 105). Here the offset described above is applied (lines 102 to 103). The test at line 99 skips over empty slots.

The buffer is released and the open is counted (lines 106 to 107).

`hdbread` (lines 111 to 119) is a character (non-buffered) read of the disk. `physio` is called to set up a buffer controller (one of the `rawbufs`) and to re-enter the driver at `hdstrat` to do the I/O. `hdstrat` must be written in anticipation of both a byte count other than 512 and the existence of a buffer address (`b_phys`) in user space.

`hdstrat` (lines 121 to 167) performs the I/O and returns. The I/O always completes before `hdstrat` returns.

The block number must be within the partition. A check is made in lines 130 to 135. There is an assumption in this driver that all I/O requests are on 512-byte block boundaries and that the byte count is a multiple of 512. The assumption is safe for file and swap I/O, but not for character I/O. Applications that disobey these constraints won't work correctly. The low 9 bits of the seek address are supplied in `b_resid` (or else ignored), and the byte count is in `b_count`.

The outer "for" loop, line 136, controls multiple block transfers. An attempt is made to do the transfer in one operation. However, the BIOS imposes constraints. Only 127 blocks may be transferred, and the I/O may not cross a physical 64K boundary. These constraints are enforced in lines 138 to 142. The low 16 bits of the physical memory address are examined at line 138. If the 64K boundary is in the middle of a block, no I/O can be done (line 139 to 140). The number of sectors transferred, `j`, is the minimum of the number left in the 64K bank and the number asked for (lines 141 to 142).

The inner "for" loop (line 143) tries 5 times to do a transfer via a BIOS call. `int13` (line 144) does a BIOS call. If all is well, `int13` returns 0 and transfer is complete (line 147). One error condition is a corrected ECC error; that error is really a warning that is taken as success (lines

148 to 152). An I/O error message is printed on the console, unless this is the first error (lines 155 to 156). The error is recovered by issuing a "reset" and the "recalibrate" command to the BIOS, causing the loop to repeat (at line 143).

If the transfer could not complete successfully after 5 tries, the I/O operation is abandoned (lines 160 to 161). Otherwise the rest of the blocks are read (the loop repeats at line 136).

The number of bytes that could not be transferred is reported to the resident in the **b\_resid** field (lines 163 to 164). If all the bytes were not transferred, a console message is printed.

The I/O is reported done at line 165.

**iotick** charges the current process 3/60ths of a second for this I/O, (line 166). This is to help out the IDRIS scheduler in case there is no system clock. This call would not be needed if the driver road-blocks while waiting for an interrupt.

**int13** (lines 179 to 203) computes the seek address given the block number and calls the BIOS interface routine, which is written in the assembler module **dos13.s**. **op** is the operation to perform; **dev** has the minor number; **nblk** is the number of 512-byte blocks to transfer; **blkno** is the block number; **offset** is the 16-bit offset within the paragraph; and **base** is the 16-bit paragraph number.

**mappart** (lines 205 to 214) returns the address of the in-memory partition table entry given the minor number.

## NAME

pchd.c code - hard disk handler

```

1:  /*  IBM PC BIOS ROM HARD DISK HANDLER
2:      *  copyright (c) 1983 by Whitesmiths, Ltd.
3:      */
4:  #include <std.h>
5:  #include <res.h>
6:  #include <bio.h>
7:
8:  /*  codes
9:      */
10: #define BIOSRS      0
11: #define BIOSRD      2
12: #define BIOSWR      3
13: #define BIOSRC      0x11
14: #define CYLBITS     0xc0
15:
16: /*  minor number encoding
17:     */
18: #define UNITSHFT     3
19: #define NHD          2
20: #define NPART        5
21: #define NDPART       4
22:
23: /*  disk tables
24:     */
25: typedef struct {
26:     UCOUNT  hd_open, hd_nspt, hd_ncyl, hd_nspc;
27: } HDTAB;
28:
29: typedef struct {
30:     UTINY pt_boot, pt_head, pt_sec, pt_cyl;
31:     UTINY pt_sysid, pt_ehead, pt_esec, pt_ecyl;
32:     ULONG pt_secoff, pt_nsec;
33: } PARTITION;
34:
35: typedef struct {
36:     UTINY br_bootstrap[0x1be];
37:     PARTITION br_part[NDPART];
38:     UTINY br_hex55, br_hexaa;
39: } BOOTREC;
40:
41: BDEVSW hdbdev {&hdopen, &hdclose, &hdstrat, &hdtab, nmhd};
42: CDEVSW hdcdev {&hdopen, &hdclose, &hdread, &hdwrite, \
43:     &nulldev, nmrhd};
44: LOCAL DEVTAB hdtab {0};
45: LOCAL HDTAB hd[NHD] {0};
46: LOCAL PARTITION hdp[NHD][NPART] {0};
47: LOCAL TEXT nmhd[] {"hdc0 hdc1 hdc2 hdc3 hdc4    hdd0 \
48:     hdd1 hdd2 hdd3 hdd4"};
49: LOCAL TEXT nmrhd[] {"rhdc0 rhdc1 rhdc2 rhdc3 rhdc4    rhdd0 rhdd1 rhdd2 \

```

```

                                rhdd3 rhdd4");
49:
50: /* close hd
51: */
52: LOCAL VOID hdclose(dev)
53:     DEV dev;
54:     {
55:     IMPORT HDTAB hd[];
56:
57:     --hd[dminor(dev) >> UNITSHFT].hd_open;
58:     }
59:
60: /* open a disk, get partition table
61: */
62: LOCAL VOID hdopen(dev)
63:     DEV dev;
64:     {
65:     IMPORT BLOCK _rootoff;
66:     IMPORT BYTES dseg;
67:     IMPORT PARTITION hdp[][];
68:     IMPORT HDTAB hd[];
69:     FAST UCOUNT hdnum = dminor(dev) >> UNITSHFT;
70:     FAST HDTAB *p = &hd[hdnum];
71:     FAST BOOTREC *pbr;
72:     BUF *pb;
73:     COUNT n, m;
74:     struct {
75:         UTINY st_sec, st_cyl, st_hdnum, st_head;
76:     } st;
77:
78:     if (NHD <= hdnum)
79:         uerror(ENXIO);
80:     else if (p->hd_open)
81:         ++p->hd_open;
82:     else if (dos13st(hdnum | 0x80, &st))
83:         uerror(ENXIO);
84:     else if (dos13(1, BIOSRD, pbr = getaddr(pb = getblk(NODEV)),
85:         1, 0, hdnum | 0x80, 0, dseg) ||
86:         pbr->br_hex55 != 0x55 || pbr->br_hexaa != 0xaa)
87:     {
88:         putfmt("bad partition table on hard disk %i\n", hdnum);
89:         brelse(pb);
90:         uerror(ENXIO);
91:     }
92:     else
93:     {
94:         p->hd_ncyl = ((st.st_sec & CYLBITS) << 2 | st.st_cyl) + 1;
95:         p->hd_nspt = (st.st_sec & ~CYLBITS);
96:         p->hd_nspc = (st.st_head + 1) * p->hd_nspt;
97:         hdp[hdnum][0].pt_nsec = (ULONG) p->hd_nspc * p->hd_ncyl;
98:         for (n = 0, m = 1; n < NDPART; ++n)
99:             if (pbr->br_part[n].pt_nsec)
100:             {
101:                 movbuf(&pbr->br_part[n], &hdp[hdnum][m], \

```

```

102:         sizeof(PARTITION));
103:         hdp[hnum][m].pt_secoff += _rootoff;
104:         hdp[hnum][m].pt_nsec -= _rootoff;
105:         ++m;
106:     }
107:     brelse(pb);
108:     ++p->hd_open;
109: }
110:
111: /* unstructured read
112: */
113: LOCAL VOID hdpread(dev)
114:     DEV dev;
115: {
116:     IMPORT VOID hdstrat();
117:
118:     physio(&hdstrat, dev, B_READ);
119: }
120:
121: /* perform I/O request
122: */
123: LOCAL VOID hdstrat(pb)
124:     FAST BUF *pb;
125: {
126:     IMPORT PARTITION *mappart();
127:     FAST BYTES i, n;
128:     UCOUNT err, j, nerr;
129:
130:     if (mappart(pb->b_dev)->pt_nsec < (n = pb->b_count \
131:         >> 9) + pb->b_blkno)
132:     {
133:         pb->b_flag |= B_ERROR;
134:         iodone(pb);
135:         return;
136:     }
137:     for (i = 0; i < n; i += j)
138:     {
139:         j = (UCOUNT)pb->b_phys;
140:         if ((127<<9) < j)
141:             break;
142:         j = 128 - ((j + 511) >> 9);
143:         j = minu(n - i, j);
144:         for (nerr = 0; nerr < 5; ++nerr)
145:             if (! (err = int13(pb->b_flag & B_READ ? BIOSRD : \
146:                 BIOSWR, pb->b_dev,
147:                 j, pb->b_blkno + i,
148:                 (BYTES)pb->b_phys & 017, (BYTES)(pb->b_phys >> 4) \
149:                     + (i << 5))))
150:                 break;
151:         else if ((err >> 8) == 0x11) /* corrected ecc */
152:         {
153:             err = 0;
154:             break;

```

```

152:         }
153:     else
154:     {
155:         if (nerr)
156:             deverr(pb, "try #i: error code %h", nerr + 1, err);
157:         int13(BIOSRS, pb->b_dev, 0, 0, 0, 0);
158:         int13(BIOSRC, pb->b_dev, 0, 0, 0, 0);
159:     }
160:     if (err)
161:         break;
162:     }
163:     if (pb->b_resid == ((n - i) << 9) + (pb->b_count & 0777))
164:         deverr(pb, "failed; %i bytes not transferred", \
            pb->b_resid);
165:     iodone(pb);
166:     iotick(3);
167:     }
168:
169: /* unstructured write
170: */
171: LOCAL VOID hdwrite(dev)
172:     DEV dev;
173:     {
174:         IMPORT VOID hdstrat();
175:
176:         physio(&hdstrat, dev, 0);
177:     }
178:
179: /* call bios to do I/O
180: */
181: LOCAL UCOUNT int13(op, dev, nblk, blkno, offset, base)
182:     UCOUNT op;
183:     DEV dev;
184:     UCOUNT nblk, blkno;
185:     BYTES offset, base;
186:     {
187:         IMPORT HDTAB hd[];
188:         IMPORT PARTITION hdp[][][], *mappart();
189:         FAST UCOUNT hdnum = dminor(dev) >> UNITSHFT;
190:         FAST HDTAB *p = &hd[hdnum];
191:         FAST UCOUNT sec;
192:         UCOUNT head, cyl;
193:         ULONG lblk;
194:
195:         lblk = blkno + mappart(dev)->pt_secoff;
196:         cyl = lblk / p->hd_nspc;
197:         sec = lblk % p->hd_nspc;
198:         head = sec / p->hd_nspt;
199:         sec = sec % p->hd_nspt;
200:         ++sec;
201:         return (dos13(nblk, op, offset, cyl >> 2 & CYLBITS | \
            sec, cyl,
            hdnum | 0x80, head, base));
202:
203:     }

```



## Appendix

pchd.c code

```
204:
205: /* map dev minor number to partition table address
206: */
207: LOCAL PARTITION *mappart(dev)
208:     DEV dev;
209:     {
210:     IMPORT PARTITION hdp[[]];
211:     FAST COUNT devm = dminor(dev);
212:
213:     return (&hdp[devm >> UNITSHFT][devm & ((1<<UNITSHFT)-1)]);
214:     }
```

## NAME

res.h - header file used by all device drivers

res.h is included in all IDRIS C source files. The ordering of include files is:

```
include <std.h>
include <res.h>
<other IDRIS include files in alphabetical order>
"local include files"
```

Pseudo types (lines 7 to 19) are:

BLOCK: holds logical block numbers  
BUF: structure of a buffer controller  
DEV: holds a device code major/minor pair  
ERROR: signed, holds an error code  
FID: file identification number (small integer)  
FVAR: structure of a file variable  
INODE: structure of an in-memory inode  
INUM: holds an inode number  
PID: holds a process id  
PROC: structure of a process list entry  
UID: holds a user id number  
USER: structure of a process table, swapped with the process image  
ZLIST: structure of a zombie process

dminor (line 36) is a macro for portably extracting the minor device number from a DEV code. dmajor (line 37) extracts the major number.

Error codes are defined in lines 46 to 78.

Scheduling priorities are defined in lines 82 to 93. These are the arguments to sleep. Processes sleeping at a negative priority cannot be interrupted. PITY and POTY are for interactive processes. PLOCK is only honored by the nice system call used for locking a process in memory. Use PBIO for sleeping on a block I/O.

Signal (software interrupt) numbers are defined in lines 97 to 114.

## NAME

res.h code - IDRIS header file

```

1:  /*  HEADER FOR IDRIS SYSTEM
2:      *  copyright (c) 1979 by Whitesmiths, Ltd.
3:      */
4:
5:  /*  basics
6:      */
7:  #define BLOCK    unsigned short
8:  #define BUF      struct buf
9:  #define DEV      unsigned short
10: #define ERROR     short
11: #define FID       short
12: #define FVAR      struct fvar
13: #define INODE     struct inode
14: #define INUM      unsigned short
15: #define PID       short
16: #define PROC      struct proc
17: #define UID       char
18: #define USER      struct user
19: #define ZLIST     struct zlist
20:
21: #define I_READ    000400
22: #define I_WRITE   000200
23: #define I_EXEC    000100
24:
25: #define NCRE      01
26: #define NDEL      02
27: #define NSYS      04
28:
29: #define NODEV     0
30: #define NOSIG     1
31: #define ROOTINO   1
32:
33: /*  the I/O device designator
34:     */
35: #define AUTOMOUNT 0x80
36: #define dminor(dev) ((dev) & BYTMASK)
37: #define dmajor(dev) (((dev) >> 8) & ~AUTOMOUNT)
38: #define isautomnt(dev) ((dev) & (AUTOMOUNT<<8))
39: struct {
40:     UTINY d_minor;
41:     UTINY d_major;
42: };
43:
44: /*  codes for u_error
45:     */
46: #define EPERM     1
47: #define ENOENT    2
48: #define ESRCH     3
49: #define EINTR     4
50: #define EIO       5

```

## res.h code

## Appendix

```
51: #define ENXIO 6
52: #define E2BIG 7
53: #define ENOEXEC 8
54: #define EBADF 9
55: #define ECHILD 10
56: #define EAGAIN 11
57: #define ENOMEM 12
58: #define EACCES 13
59: #define ENOTBLK 15
60: #define EBUSY 16
61: #define EEXIST 17
62: #define EXDEV 18
63: #define ENODEV 19
64: #define ENOTDIR 20
65: #define EISDIR 21
66: #define EINVAL 22
67: #define ENFILE 23
68: #define EMFILE 24
69: #define ENOTTY 25
70: #define ETXTBSY 26
71: #define EFBIG 27
72: #define ENOSPC 28
73: #define ESPIPE 29
74: #define EROFS 30
75: #define EMLINK 31
76: #define EPIPE 32
77: #define ENOSYS 100
78: #define EFAULT 106
79:
80: /* priorities
81: */
82: #define PLOCK -20
83: #define PSWAP -3
84: #define PIIIO -2
85: #define PBIO -1
86: #define PITY 0
87: #define POTY 0
88: #define PNORM 4
89: #define PINIT 4
90: #define PWAIT 10
91: #define PPIPE 11
92: #define PSLEEP 20
93: #define PLOW 20
94:
95: /* signals
96: */
97: #define SIGHUP 1
98: #define SIGINT 2
99: #define SIGQUIT 3
100: #define SIGKILL 4
101: #define SIGTRAP 5
102: #define SIGTSTP 6 /* IOT */
103: #define SIGTERM 7 /* EMT */
104: #define SIGABRT 8
```

## Appendix

res.h code

```
105: #define SIGKIL 9
106: #define SIGBUS 10
107: #define SIGSEG 11
108: #define SIGSYS 12
109: #define SIGPIPE 13
110: #define SIGALRM 14
111: #define SIGTERM 15
112: #define SIGSVC 16
113: #define SIGUSR 17
114: #define NSIG 18 /* must be last sig + 1 */
```